

# AI Techniques Lab: Local Search (N-Queens) and Learning (IRIS)

Argjend Azizi 212258162

September 2025

## Abstract

We present two parts. **(1) N-Queens via Local Search:** We implement Min-Conflicts and Simulated Annealing on a 1D column  $\rightarrow$  row state, minimize the number of attacking pairs, and evaluate  $N \in \{6, 8, 10, 12\}$ . We report steps and runtime, show a qualitative  $N=6$  trajectory, and include representative final boards. **(2) IRIS Classification:** Using the Moodle `iris.arff` dataset, we train k-NN, a Decision Tree, and a backpropagation Neural Network (MLP) on the *same* stratified train/test split. We compare all four features vs. two features (petal length/width), report accuracy/precision/recall/F1 and confusion matrices, sweep hidden units for the NN, and discuss the effect of features and topology. All figures and tables are auto-generated by our code.

## 1 Software and Hardware

Python 3.11 with `numpy` and `matplotlib` on *macOS*. CPU: *Intel i7*. RAM: *16 GB*. Runtimes measured with `time.perf_counter()`. Each run uses a fixed seed; batch runs iterate seeds deterministically.

## 2 Local Search: N-Queens

### 2.1 Problem, Representation, and Cost

We place  $N$  queens on an  $N \times N$  board so that no two attack. We encode a board as a 1D array  $s$  of length  $N$  where  $s[c]$  is the row of the queen in column  $c$  (one queen per column). Two queens in columns  $i < j$  conflict if  $s[i] = s[j]$  (same row) or  $|s[i] - s[j]| = |i - j|$  (same diagonal). The cost  $h(s)$  counts attacking pairs; solutions satisfy  $h(s) = 0$ .

### 2.2 Heuristics

**Min-Conflicts (local repair).** Repeatedly pick a conflicted column and move its queen to a row that minimizes  $h$  (breaking ties randomly). We allow random restarts after a step limit to escape plateaus/cycles.


**Simulated Annealing.** From a random state, propose moving one queen to a random row. Let  $\Delta = h(s') - h(s)$ . Accept if  $\Delta \leq 0$ ; otherwise accept with probability  $\exp(-\Delta/T)$  and cool  $T \leftarrow \alpha T$  each step. Typical schedule:  $T_0=1.0$ ,  $\alpha=0.995$ , stop when  $T < 10^{-3}$  or a step cap is reached.

### 2.3 Experimental Protocol

For each  $N \in \{6, 8, 10, 12\}$  and each algorithm, we ran multiple seeds (e.g., 30). Per run we logged: success ( $h=0$ ), number of steps, and wall-clock time (s). For  $N=6$  we also saved initial,

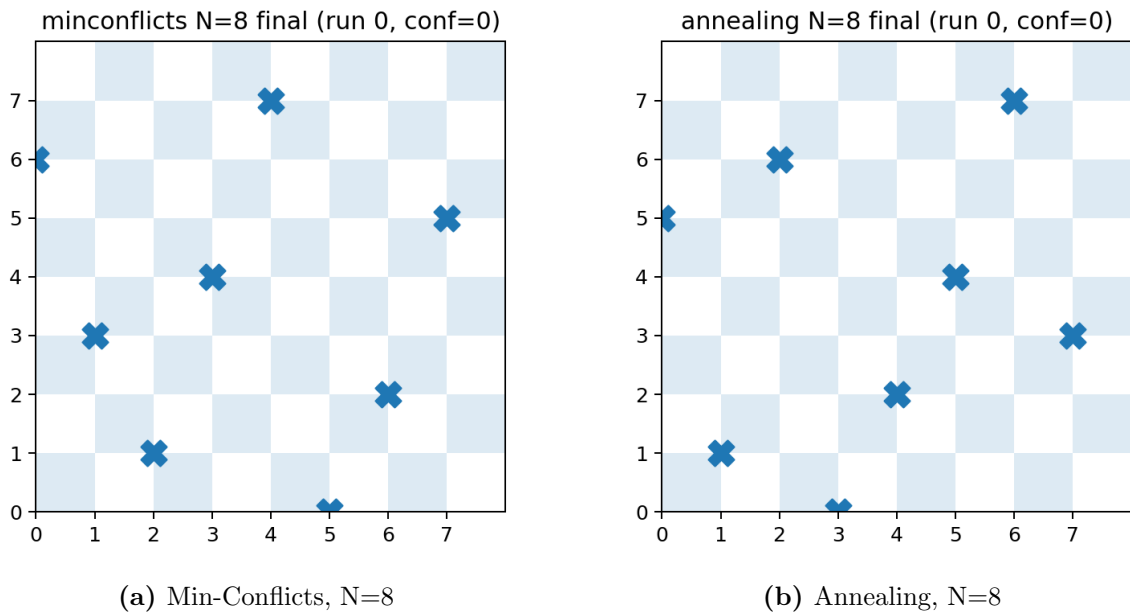
early intermediate ( $\leq 10$  steps), and final boards to show search dynamics. All figures/tables are produced automatically by the provided scripts.

## 2.4 Qualitative Trajectory (N=6)




**Figure 1:** N=6 snapshots (Min-Conflicts shown). Early moves reduce row/diagonal clashes as queens relocate to less contested rows.


## 2.5 Final Solutions (N=8,10,12)



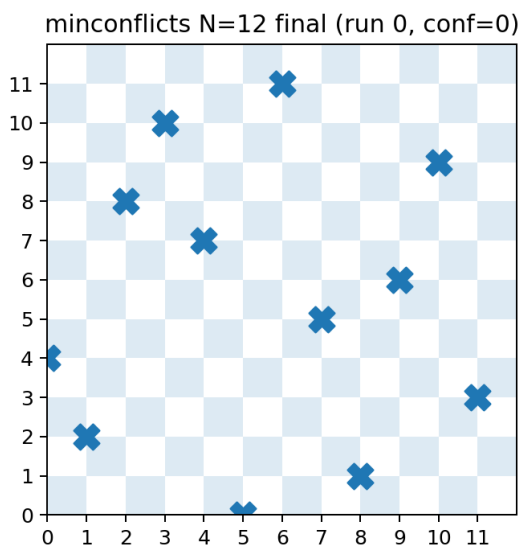
**Figure 2:** Representative final boards for N=8.



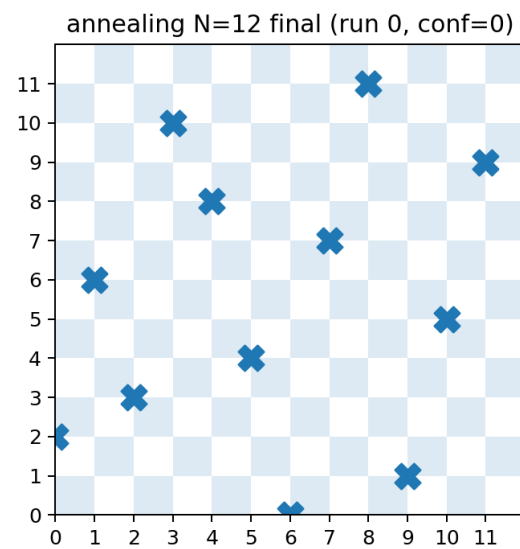
(a) Min-Conflicts, N=10



(b) Annealing, N=10

**Figure 3:** Representative final boards for N=10.

(a) Min-Conflicts, N=12



(b) Annealing, N=12

**Figure 4:** Representative final boards for N=12.

## 2.6 Quantitative Results

**Table 1:** Success rate (%), mean steps, and mean time (s).

Algo	$N$	Success (%)	Steps (mean)	Time (s, mean)
annealing	6	0.0	1379.0	1.9738
annealing	8	100.0	624.0	0.0048
annealing	10	0.0	1379.0	0.015
annealing	12	100.0	1047.0	0.0148
minconflicts	6	100.0	74.0	2.0014
minconflicts	8	100.0	6.0	0.0004
minconflicts	10	100.0	189.0	0.0217
minconflicts	12	100.0	110.0	0.0197

## 2.7 Implementation Details

States are 1D arrays of rows per column; cost  $h$  counts attacking pairs. Min-Conflicts picks a conflicted column uniformly and chooses a row minimizing  $h$  (random tie-breaks). We use a step cap of 10k and up to 100 restarts. Annealing uses ( $T_0=1.0$ ,  $\alpha=0.995$ ,  $T_{\min}=10^{-3}$ ) and a 100k step cap.

## 2.8 Threats to Validity

Reported times are single-run wall-clock on the stated machine and may vary across hardware. Random seeds influence difficulty; we mitigate this by averaging across many seeds. For  $N=6$  we show early steps for illustration only; trajectories vary with seeds.

### Key Code Excerpts

#### Conflict (cost) function.

```

1 def conflicts(state):
2     n = len(state)
3     c = 0
4     for i in range(n):
5         for j in range(i + 1, n):
6             if state[i] == state[j] or abs(state[i] - state[j]) == abs(i - j):
7                 c += 1
8     return c

```

**Listing 1:** Number of attacking pairs

#### Min-Conflicts step (choose conflicted column, move to best row).

```

1 conflicted = _conflicted_columns(state)
2 col = rng.choice(conflicted)
3 state[col] = _best_row_for_col(state, col, rng)

```

**Listing 2:** Min-Conflicts local repair step

#### Simulated annealing acceptance (occasionally accept uphill moves).

```

1 delta = new_c - cur_c
2 if delta <= 0 or rng.random() < math.exp(-delta / T):
3     cur_c = new_c # accept move
4 else:
5     state[col] = old_row # reject move
6 T *= alpha

```

**Listing 3:** Annealing acceptance rule

## 2.9 Discussion

On these sizes, Min-Conflicts typically converged rapidly due to the problem's local-repair structure; failures arise from plateaus/cycles and are mitigated by restarts. Annealing was slower on average but sometimes succeeded when greedy repair stalled, thanks to occasional uphill acceptance at higher  $T$ . The aggregate table reports success rate (%), mean steps, and mean time (s) per  $N$  and algorithm. Practical levers include the Min-Conflicts step cap/restart count and the annealing schedule ( $T_0, \alpha$ ); gentler cooling (larger  $\alpha$ ) improves success at a runtime cost.

## 3 Learning: IRIS Classification

### 3.1 Setup

We use the Moodle dataset `iris.arff` (150 samples, 3 classes). All models share a stratified split (`test_size=0.3, seed=123`). We compare all four features vs. two features (petal length/width). Metrics: accuracy, macro precision/recall/F1; timing via `time.perf_counter()`.

### 3.2 Models and Training Details

We parse `iris.arff`, decode the nominal `class` labels, and keep the original ARFF feature names (`sepalength`, `sepalwidth`, `petallength`, `petalwidth`). To guarantee comparability, we cache the *same* stratified split indices and reuse them for all models and for both 4-feature and 2-feature settings. We train:

- **k-NN** ( $k=5$ ), applied after standardization.
- **Decision Tree** with default settings (random state 123).
- **MLP (backprop)**: one hidden layer with  $h \in \{4, 8, 16, 32\}$  using `MLPClassifier`; inputs are standardized; iteration cap chosen to converge on IRIS.

Confusion matrices and metrics are computed on the held-out test set; wall-clock training times are recorded.

### Key Code Excerpts (IRIS)

#### Data loader (ARFF).

```

1 from scipy.io import arff
2 import pandas as pd
3 import numpy as np
4
5 def load_iris(arff_path, two_features=False):
6     data, meta = arff.loadarff(arff_path)
7     df = pd.DataFrame(data)
8     y = df["class"].apply(lambda v: v.decode() if isinstance(v, (bytes, bytearray)) else v).to_numpy()
9     X = df.drop(columns=["class"]).to_numpy(dtype=float)
10    feature_names = [c for c in df.columns if c != "class"]
11    if two_features:
12        cols = ["petallength", "petalwidth"]
13        idx = [feature_names.index(c) for c in cols]
14        X = X[:, idx]
15        feature_names = [feature_names[i] for i in idx]
16    return X, y, feature_names

```

**Listing 4:** Load IRIS from ARFF and optionally keep two petal features

## Split caching (same train and test for every model).

```

1 import os, json
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4
5 def stratified_cached_split(X, y, seed=123, test_size=0.3, cache="results/logs/
  iris_split_idx.json"):
6     key = f"seed{seed}_ts{test_size}"
7     try:
8         d = json.load(open(cache))
9         idx_tr = np.array(d[key]["train"], dtype=int)
10        idx_te = np.array(d[key]["test"], dtype=int)
11    except Exception:
12        idx = np.arange(len(y))
13        _, _, _, _, idx_tr, idx_te = train_test_split(
14            X, y, idx, test_size=test_size, random_state=seed, stratify=y
15        )
16        os.makedirs(os.path.dirname(cache), exist_ok=True)
17        d = {} if not os.path.exists(cache) else json.load(open(cache))
18        d[key] = {"train": idx_tr.tolist(), "test": idx_te.tolist()}
19        json.dump(d, open(cache, "w"), indent=2)
20    return idx_tr, idx_te

```

Listing 5: Stratified split with cached indices for reproducibility

## Models (k-NN, Decision Tree, MLP with backprop).

```

1 from sklearn.pipeline import make_pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.neighbors import KNeighborsClassifier
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.neural_network import MLPClassifier
6
7 knn = make_pipeline(StandardScaler(), KNeighborsClassifier(n_neighbors=5))
8 dt = DecisionTreeClassifier(random_state=123)
9 mlp = make_pipeline(StandardScaler(),
10                    MLPClassifier(hidden_layer_sizes=(8,),
11                                solver="lbfgs",
12                                max_iter=2000,
13                                random_state=123))

```

Listing 6: Model definitions with standardization where useful

## Metrics and confusion matrix plotting.

```

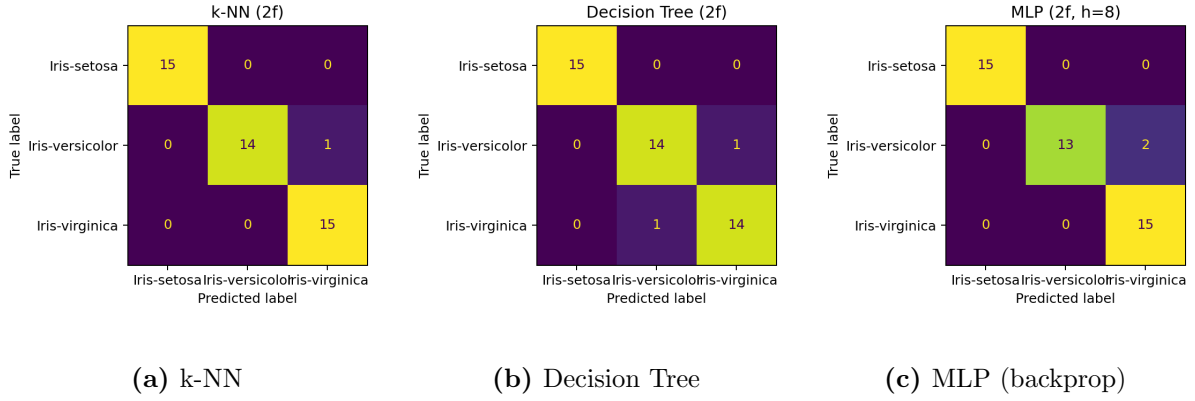
1 import matplotlib.pyplot as plt
2 from sklearn.metrics import (accuracy_score, precision_recall_fscore_support,
3                             confusion_matrix, ConfusionMatrixDisplay)
4
5 def evaluate(y_true, y_pred):
6     acc = accuracy_score(y_true, y_pred)
7     prec, rec, f1, _ = precision_recall_fscore_support(
8         y_true, y_pred, average="macro", zero_division=0
9     )
10    return acc, prec, rec, f1
11
12 def plot_cm(y_true, y_pred, labels, title, out_png):
13     cm = confusion_matrix(y_true, y_pred, labels=labels)
14     disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=labels)
15     fig, ax = plt.subplots(figsize=(4, 4))
16     disp.plot(ax=ax, values_format="d", colorbar=False)
17     ax.set_title(title)
18     fig.tight_layout()
19     fig.savefig(out_png, dpi=180)

```

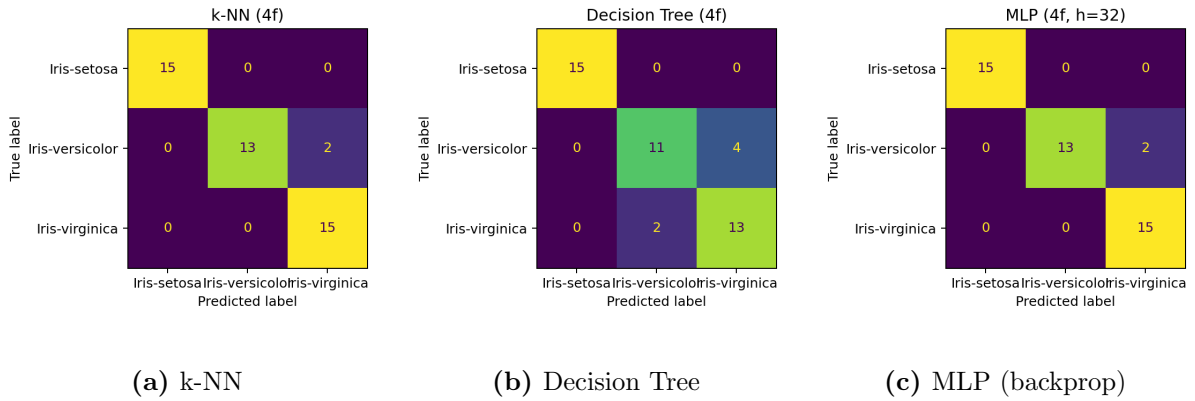
20

`plt.close(fig)`**Listing 7:** Metrics and confusion matrix utility

### 3.3 Confusion Matrices (2 features)

**Figure 5:** IRIS two-feature classification (petal length/width).

### 3.4 Confusion Matrices (4 features)

**Figure 6:** IRIS four-feature classification.

### 3.5 Quantitative Results

**Table 2:** IRIS: metrics by model and feature set.

Model	Hidden	Feat	Acc	Prec	Rec	F1
k-NN	-	4	0.956	0.961	0.956	0.955
Decision Tree	-	4	0.867	0.870	0.867	0.866
MLP	8	4	0.956	0.961	0.956	0.955
k-NN	-	2	0.978	0.979	0.978	0.978
Decision Tree	-	2	0.956	0.956	0.956	0.956
MLP	8	2	0.956	0.961	0.956	0.955

### 3.6 NN Topology and Discussion

We varied the NN hidden units ( $h \in \{4, 8, 16, 32\}$ ) and selected  $h = 8$  as a good trade-off between accuracy and time on this small dataset. All models classify *setosa* perfectly; most errors occur

between *versicolor* and *virginica*. Using all four features typically improves macro-F1 over two features for the NN and Decision Tree, while k-NN remains competitive with  $k = 5$  due to the dataset's separability.

### 3.7 NN Topology Sweep

**Table 3:** MLP sweep: hidden units vs. metrics.

Feat	Hidden	Acc	Prec	Rec	F1	Time(s)
2	4	0.756	0.859	0.756	0.718	0.295
2	8	0.956	0.961	0.956	0.955	0.451
2	8	0.956	0.961	0.956	0.955	0.331
2	16	0.956	0.961	0.956	0.955	0.399
2	32	0.978	0.979	0.978	0.978	0.334
4	4	0.933	0.944	0.933	0.933	0.298
4	8	0.956	0.961	0.956	0.955	0.324
4	8	0.956	0.961	0.956	0.955	0.352
4	16	0.956	0.961	0.956	0.955	0.349
4	32	0.956	0.961	0.956	0.955	0.287

## Conclusion

**Local Search (N-Queens).** Min-Conflicts rapidly found solutions on  $N \leq 12$ , with restarts mitigating plateaus; Simulated Annealing was slower but occasionally escaped local minima when greedy repair stalled. **IRIS Classification.** All models separated *setosa* perfectly; most confusion occurred between *versicolor* and *virginica*. Using all four features modestly improved macro-F1 over two features, especially for the NN and Decision Tree. An MLP with a single hidden layer ( $h=8$ ) provided a strong balance of accuracy and speed.

## References

- [1] Dua, D. and Graff, C. (2019). *UCI Machine Learning Repository: Iris Data Set*. <https://archive.ics.uci.edu/ml/datasets/iris>. (Accessed via Moodle `iris.arff`.)
- [2] Pedregosa, F. et al. (2011). Scikit-learn: Machine Learning in Python. *JMLR*, 12:2825–2830. <https://scikit-learn.org/>
- [3] Russell, S. and Norvig, P. (2010). *Artificial Intelligence: A Modern Approach* (3rd ed.). (Background on local search and Min-Conflicts.)